

# Autoscaling in Kubernetes

Abdullah Al Hasib

[abdullah.alhasib@kratosdefense.com](mailto:abdullah.alhasib@kratosdefense.com)

09-02-2022



① Kubernetes Recap

② Autoscaling

③ HPA

④ VPA

⑤ CA

⑥ Conclusion

## ① Kubernetes Recap

## ② Autoscaling

## ③ HPA


## ④ VPA

## ⑤ CA

## ⑥ Conclusion

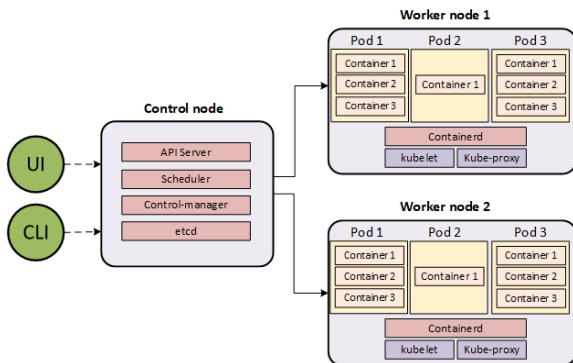
## A Quick Recap

### What is kubernetes:



# A Quick Recap

## Kubernetes architecture:



① Kubernetes Recap

② Autoscaling

③ HPA

④ VPA

⑤ CA

⑥ Conclusion

# Autoscaling

## Autoscaling:

- automatically adjusts computational resource usage according to the load.

## Why do we need it:

- to cope with the demand
- to reduce cost
- to reduce power consumption

# Autoscaling Types

The Kubernetes autoscaling mechanism uses two layers:

- Pod-based scaling - supported by
  - Horizontal Pod Autoscaler (HPA)
  - Vertical Pod Autoscaler (VPA)
- Node-based scaling - supported by
  - Cluster Autoscaler (CA)

	Pods	Nodes
Horizontal	# of pods	# of nodes
Vertical	resources of a pod	resources of a node



1 Kubernetes Recap

2 Autoscaling

3 HPA


4 VPA

5 CA

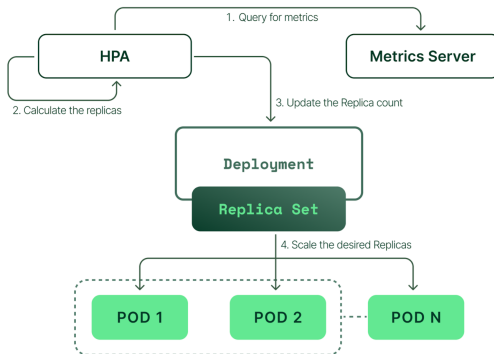
6 Conclusion

# Horizontal Pod Autoscaling

HPA increases or decreases the number of pods in a replication controller, deployment, replica set etc. based on CPU utilization



# How Does HPA Work?



```
desiredReplicas = ceil[currentReplicas * ( currentMetricValue / desiredMetricValue )]
```

# Metrics Types

Types of metrics APIs:

- 1 Resource Metrics (`metrics.k8s.io`)
  - predefined resource usage metrics of Pods and Nodes
  - can be expressed as raw value of percentage
  - example: CPU, Memory

# Metrics Types

Types of metrics APIs:

- 1 Resource Metrics (`metrics.k8s.io`)
  - predefined resource usage metrics of Pods and Nodes
  - can be expressed as raw value of percentage
  - example: CPU, Memory
- 2 Custom Metrics (`custom.metrics.k8s.io`)
  - custom metrics associated with a Kubernetes object
  - example: `rate_of_client_requests`

# Metrics Types

Types of metrics APIs:

- ① Resource Metrics (`metrics.k8s.io`)
  - predefined resource usage metrics of Pods and Nodes
  - can be expressed as raw value of percentage
  - example: CPU, Memory
- ② Custom Metrics (`custom.metrics.k8s.io`)
  - custom metrics associated with a Kubernetes object
  - example: `rate_of_client_requests`
- ③ External Metrics (`external.metrics.k8s.io`)
  - custom metrics not associated with a Kubernetes object

## Customizing Scaling Behavior

- ① `scaleup`: control scaling behavior while scaling up
  - `stabilizationWindowSeconds`: (default:0)
  - `selectPolicy`: can be Min, Max, Disabled (default:Max)
  - `policies`
    - `type`: Pods or Percent
    - `periodSeconds`: (default: 60)
    - `value`

## Customizing Scaling Behavior

- ① **scaleup**: control scaling behavior while scaling up
  - `stabilizationWindowSeconds`: (default:0)
  - `selectPolicy`: can be Min, Max, Disabled (default:Max)
  - `policies`
    - `type`: Pods or Percent
    - `periodSeconds`: (default: 60)
    - `value`
- ② **scaledown**: control scaling behavior while scaling down
  - `stabilizationWindowSeconds`: (default: 300)
  - `selectPolicy`: can be Min, Max, Disabled (default:Max)
  - `policies`
    - `type`: Pods or Percent
    - `periodSeconds`: (default: 60)
    - `value`



# User Stories

- 1 Scale up as fast as possible, scale down as usual

```
1 behavior:
2   scaleUp:
3     policies:
4       - type: Percent
5         value: 900%
```

# User Stories

- ① Scale up as fast as possible, scale down as usual

```
1 behavior:
2   scaleUp:
3     policies:
4       - type: Percent
5         value: 900%
```

- ② Scale up as usual, do not scale down

```
1 behavior:
2   scaleDown:
3     selectPolicy: Disabled
```

# User Stories

## ④ Scale Up As Fast As Possible, Scale Down Very Gradually

```
1  behavior:
2    scaleUp:
3      policies:
4        - type: Percent
5          value: 900%
6    scaleDown:
7      policies:
8        - type: Pods
9          value: 1
10     periodSeconds: 600
```

## User Stories (Cont'd)

### 5 Stabilization before scaling down

```
1 behavior:
2   scaleDown:
3     stabilizationWindowSeconds: 600
4   policies:
5     - type: Pods
6       value: 5
```

## HPA Demo

# HPA Demo

1 Kubernetes Recap

2 Autoscaling

3 HPA


4 VPA

5 CA

6 Conclusion

# Vertical Pod Autoscaling

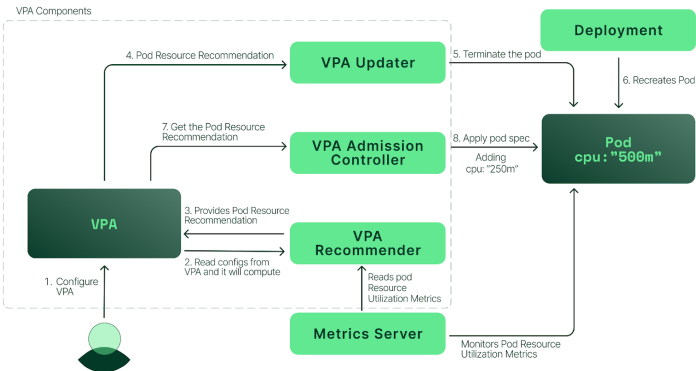
VPA adjusts the resource requests and limits of containers in the cluster.



Resource configuration types:

- requests: define the minimum amount of resources that containers need
- limits: define the maximum amount of resources that a given container can consume

# How Does VPA Work?





# VPA Concepts

## Types Operation Modes:

- Off: provides recommendations only
- Initial: only assigns resource requests on pod creation
- Recreate: assigns resource requests on pod creation and updates them on existing pods by evicting them. [should be used rarely]
- Auto: currently does the same as Recreate. [may cause service downtime]

## VPA Demo

# VPA Demo

1 Kubernetes Recap

2 Autoscaling

3 HPA

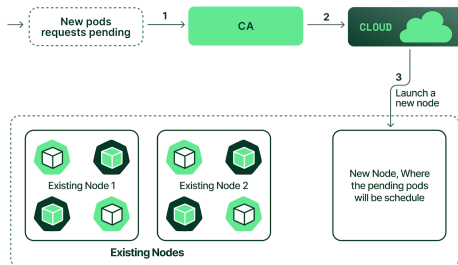
4 VPA

5 CA

6 Conclusion

# Cluster Autoscaling

CA adjusts the number of nodes in the cluster when pods fail to schedule or when nodes are underutilized



① Kubernetes Recap

② Autoscaling

③ HPA

④ VPA

⑤ CA

⑥ Conclusion

## Conclusion

- Install a metric server
- Do not mix HPA with VPA
- Define pod requests and limits
- Resource requests should be close to the average usage of the pods
- Increase CPU limits for slow starting applications

- [1] Minikube Handbook
- [2] Horizontal Pod Autoscaling
- [3] Vertical Pod Autoscaling
- [4] Scaling Kubernetes Clusters
- [5] KEDA: Kubernetes Event-Driven Autoscaling

*Thanks!*